qAI37

INTRODUCTION DECK

The Neutral Route to Post-Silicon AI

An AI infrastructure company building the vendor-agnostic software access point between conventional AI and neutral-atom systems.

CONFIDENTIAL · DOOR ONE · PRE-NDA

Version 1.2 · For discussion purposes only

Ted Stockwell | teds@qai37.com
Michelle Holtmann | michelle@qai37.com

See legal notice on Slide 14.

The Problem Is Not a Lack of Chips

The present constraint is architectural: long-context inference stops being interactive.

55–137 s

1M tokens

Modeled, per turn · ten-turn session basis

3–8 min

5M tokens

Modeled, per turn · ten-turn session basis

12–30 min

10M tokens

Modeled, per turn · ten-turn session basis

More hardware shrinks the wait. It does not remove the wall.

- Transformer attention repeatedly works against the expanding history — the whole context is re-read before every answer.
- Memory movement and power rise with context and concurrency.
- Full-codebase agents, enterprise knowledge bases, and multi-hour video become non-interactive.

All timing figures are engineering projections from published 8×H100-class specifications; assumptions and derivation are available in controlled diligence.

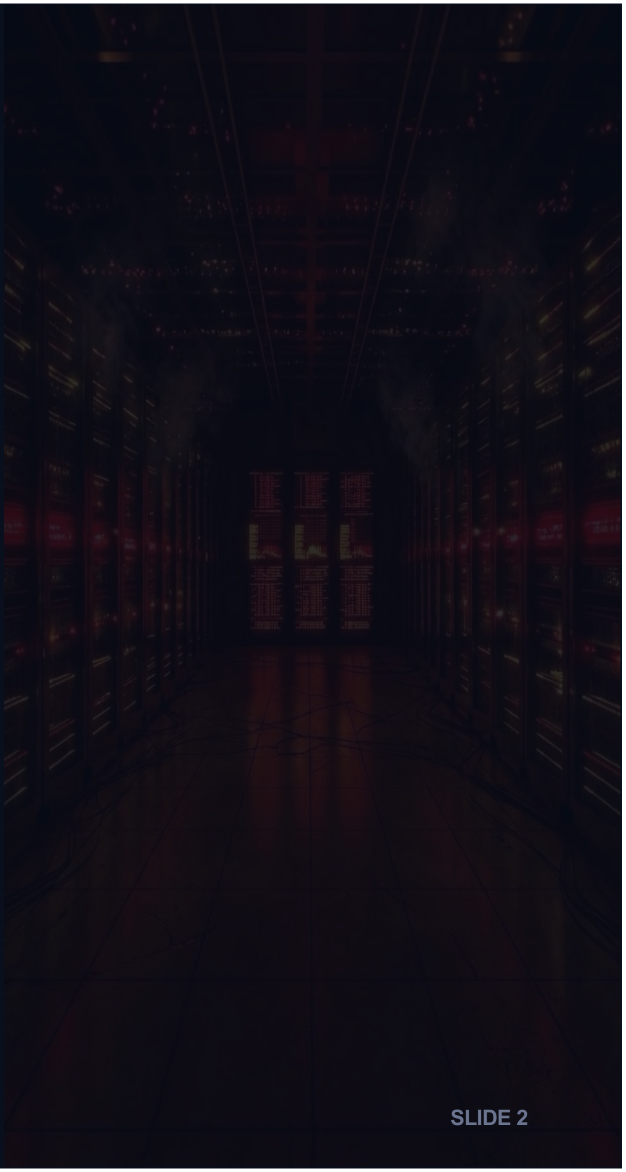

This Is an AI Infrastructure Company

Quantum describes the processor class beneath the software layer. It does not define the company category.

1 What stays conventional

The customer’s model, application, data, orchestration, security boundary, and classical fallback remain in place.

2 What qAI37 changes

A selected computational path moves beneath the API to a target backend. qAI37 handles qualification, translation, portability, and return.

3 What qAI37 is not

Not another quantum machine. Not a replacement for the data center. Not a bet that one hardware vendor wins. Not a bet against silicon.

The product is the software access point between today’s AI stack and an additional processor class.

The Canal, Not Another Ship

The Suez Canal favored no carrier and no hull; it shortened every voyage; all of the traffic passed through it. qAI37 builds the neutral passage between the installed AI estate and the neutral-atom landscape.

CLASSICAL AI

- Applications
- Models and weights
- Data and security
- Orchestration and fallback

The installed estate stays.

qAI37

THE NEUTRAL ROUTE

- Intercept
- Qualify
- Translate
- Return

One passage. Multiple destinations.

NEUTRAL-ATOM LANDSCAPE

QuEra

Pasqal

Atom Computing

Infleqtion

Oratomic

The category landscape, scaling on its own capital.

Canal framing from Sethu Raman's August 2026 technical review, paraphrased. The position is being built, not already held. Vendor names identify the category landscape — target systems, not partnerships, endorsements, or access agreements; two have working emulation paths today (Slide 5).

What Has Been Shown

A standard AI application sent one selected layer through qAI37 and received an inference result; the surrounding application remained classical.

1 Path 1 — reviewed live by Sethu Raman

A classically trained model kept its application boundary while one selected layer was routed through qAI37 — request in, real-time compilation into the machine's own control instructions, inference result out.

2 Path 2 — demonstrated separately

A second vendor-targeted neutral-atom path also ran live, end-to-end in emulation, without an application rewrite. The vendors' own physics-accurate emulators were used on both paths.

SHOWN

- Two vendor-targeted neutral-atom paths in emulation.
- Both demonstrated live; one reviewed by Sethu Raman.
- Request in, vendor-targeted execution path, result out.
- No application rewrite.

NOT YET SHOWN — EACH A FUNDED GATE

- Memory carried across turns — the scaling claim (M2).
- The automated pipeline; today's path is orchestrated by hand (M1, Proof Point Zero).
- Physical-QPU execution (M3); output quality, economics, customer proof (M4).

The first question a careful reader asks — did it remember? — has an honest answer: not yet; that is the next gate. What has been shown is that inference runs on this class of machine at all.

The Advantage Is Memory, Not Faster Compute

A machine that holds state carries context forward instead of re-reading the whole history before every answer.

C The instinct

“Quantum means a faster computer.” For this purpose it does not. Silicon is stateless: today’s AI re-reads everything it has been given before every answer, and the more it knows, the slower and costlier every answer becomes.

S The stateful path

Superposition, for this purpose, is the ability to hold state. A memory-native recurrent model — a stateful recurrent inference path — keeps working memory resident in the machine and carries it forward; only small differences cross the boundary.

P Why the field said no

Transformers on quantum processors fail, for a structural reason: the transformer moves enormous memory at enormous speed, and no channel of that bandwidth exists into a quantum machine. The field is right about the transformer and wrong about the category.

Evidence boundary: the route is demonstrated; persistent memory across turns, long-context scaling, output parity, and physical-QPU performance are not.

Published foundation: Eric R. Anschuetz and Xun Gao, Quantum 10, 1976 (2026) — trainable recurrent quantum memory separations. qAI37 does not claim the theorem; it builds the access and execution stack that makes stateful paths usable. The idea is decades old and proven; the hardware that holds state is new.

What the Product Does

The interface stays familiar; the execution path changes beneath it.

1 Intercept

Capture a supported inference call through the existing application interface.

2 Qualify

Determine whether the selected operation and target backend are eligible.

3 Translate

Map the operation into inspectable, vendor-targeted instructions.

4 Return

Bring the result back into the conventional application and fallback path.

Vendor Portability: one application integration, multiple target formats, classical fallback intact. Runtime selection among live destinations is a later operational layer.

Why Neutrality Is the Product

A vendor SDK is a necessary port into one machine. It is not a neutral route across all of them.

V The vendor’s rational position

Optimize for its own machine, its own roadmap, and its own lock-in. Supporting rivals on equal terms weakens its economics.

R qAI37’s independent position

Qualify destinations, preserve customer choice, and route workloads without requiring qAI37 to own or finance the hardware.

C The customer outcome

One software access point, multiple qualified destinations over time, and a conventional AI estate that remains under operator control.

Each qualified backend increases choice for workload owners; each qualified workload increases the value of the route to vendors.

The Moat Is Infrastructure, Not Code Alone

The company must remain defensible if software gets cheaper and hardware commoditizes.

D Operating corpus

Machine-specific calibration, qualification, failure handling, and execution history can accumulate with every run. A code generator cannot recreate a record it never observed.

N Two-sided network

More supported backends increase value to workload owners. More workloads increase value to hardware vendors. The route compounds in both directions.

I Embedded incumbency

Once the access layer sits inside a customer workflow and a vendor integration, replacement becomes a multi-quarter infrastructure decision.

The pressure test: if hardware proliferates, the route gains destinations; if software becomes easier to build, the operating record and installed position matter more.

The Hardware Is Being Funded. The Evidence Path Is Bounded.

Independent capital validates the substrate; qAI37 advances through explicit proof gates.

G Google

Entered neutral-atom quantum computing in March 2026, alongside its superconducting program.

A Atom Computing

$100M Series C plus a planned $100M Department of Commerce Letter of Intent, June 2026.

O Oratomic

Public launch March 30, 2026; $300M Series A on July 7, 2026.

TODAY	M1	M2	M3	M4
	~Month 9	~Month 12	~Month 18	~Month 24
Two paths	**Software**	**Scaling**	**Hardware**	**Customer**
End-to-end in emulation.	Proof Point Zero: inspectable mechanism; no qubits.	Target workload under emulation.	Same workload on two QPU backends.	Output, integration, reliability, economics.

The program funds evidence, not speculative hardware capacity. Public events validate category timing, not qAI37 performance.

Independent Scrutiny and an Execution Team

Quantum depth in two chairs, the HAL and translation architecture in a third, and an independent evaluator beside them.

INDEPENDENT TECHNICAL REVIEW

Sethu Raman

Former Distinguished Engineer at both Apple and Microsoft

- January–March 2026: three months of independent technical diligence within an institutional venture evaluation.
- Unpaid, with no equity at evaluation.
- Subsequently formalized as Technical Advisor.
- Available for direct technical diligence.
- Reviewed the August 2026 live demonstration (Slide 5).

Employer names appear as text only. Statuses are precise on request.

Ted Stockwell

Founder & CEO

Former GM, Microsoft Online Services Division; built Bing as a Platform.

Michelle Holtmann

President & CSO

Former GM, Microsoft Enterprise Infrastructure; platforms serving billions of users.

Rupesh Srivastava

Quantum Advisor

Quantum ecosystem. PhD Physics, Royal Holloway, University of London; five years developing the UK quantum-computing ecosystem at Oxford's Department of Physics; CQO, Entangled Positions.

Steve Jahnke

CTO / Principal Architect

Owns the HAL and translation architecture. 30 years at Intel, Altera, and Texas Instruments.

Vincent E. Elfving

Chief Quantum Advisor

Quantum algorithms. Former Head of Algorithms, Pasqal — led 40+ researchers on neutral-atom AI workflows. Google Quantum AI alumnus.

Laverne Masaki

Chief People Officer

Former Microsoft and Google executive recruiter; built and led organizations across MSR and OSD.

Quantum-algorithm and quantum-ecosystem depth (Elfving, Srivastava) sits beside the HAL and translation architecture (Jahnke); the evaluator sits beside the team, not inside it.

The Thesis in Three Sentences

The understanding sequence, reduced to the shortest supportable form.

1. **You're building a neural network that runs on quantum hardware.**

   Not a faster chip for today's AI — a different kind of AI on a different kind of machine. Today that route is demonstrated in emulation, not on physical hardware.

2. **It remembers — so it's better suited to inference. It works the way we do.**

   The machine holds context the way a person does, instead of re-reading everything before every answer. Persistent memory across turns is the next gate, not a claim.

3. **A breakthrough theory that was proven but never applicable is now applicable.**

   The idea is decades old and sound; the hardware that can hold state finally exists. Old proof, new hardware, and an unoccupied access-point position.

These are company statements, not quotations from a private conversation.

The Opening

1 · The software boundary works across two vendor-targeted paths in emulation.

2 · The neutral-atom hardware category is being institutionally capitalized.

3 · The independent access-point position above that category remains open.

Door Two under NDA: architecture · technical Q&A · Proof Point Zero specification · program detail · the demonstration, shown live with the engineering team

qAI37 is raising a Seed/Series A four-milestone program. We are open to a clean equity structure with a lead investor. Terms to be discussed.

Ted Stockwell | teds@qai37.com   **Michelle Holtmann | michelle@qai37.com**

CONFIDENTIAL · DOOR ONE · PRE-NDA

Legal Notice and Evidence Discipline

Door One is a pre-NDA discussion room. This notice governs the deck and its companion.

Confidential pre-NDA material

Provided through controlled DocSend access for evaluation. These materials contain confidential and proprietary information of qAI37, Inc.; do not copy, forward, or redistribute without qAI37's prior written consent.

No offer or solicitation

These materials are informational and do not constitute an offer to sell, a solicitation of an offer to buy, or a recommendation to purchase any securities. Any offer will be made solely through definitive investment documentation executed by authorized officers of the company.

Forward-looking statements

Milestones, timelines — including Proof Point Zero and subsequent validation phases — market development, and financing outcomes are forward-looking, reflect management expectations as of August 2026, and are subject to risks and execution uncertainty. qAI37 undertakes no obligation to update them.

Modeled and third-party information

Unless expressly labeled measured, performance, latency, efficiency, energy, and scaling figures are engineering estimates from published results, vendor specifications, internal modeling, and physics-accurate emulation. Third-party names, statements, and events do not imply endorsement, partnership, or a commercial relationship.

Evidence classes remain separate

Emulator evidence is not physical-QPU performance. A demonstrated route is not formal Proof Point Zero acceptance, output parity, customer proof, or commercial readiness. Each later class is a funded, checkable gate.

Questions and diligence

Specific evidence, assumptions, architecture, and program details are available through Door Two under NDA, with direct access to the team and the independent evaluator.

 | Ted Stockwell — teds@qai37.com | Michelle Holtmann — michelle@qai37.com